

sklearn.svm.SVR

```
class sklearn.svm.SVR(*, kernel='rbf', degree=3, gamma='scale', coef0=0.0, tol=0.001, C=1.0, epsilon=0.1, shrinking=True, cache_size=200, verbose=False, max_iter=-1)
```

Epsilon-Support Vector Regression.

The free parameters in the model are C and epsilon.

The implementation is based on libsvm. The fit time complexity is more than quadratic with the number of samples which makes it hard to scale to datasets with more than a couple of 10000 samples. For large datasets consider using [LinearSVR](#) or [SGDRegressor](#) instead, possibly after a [Nystroem](#) transformer.

Parameters:	kernel : {'linear', 'poly', 'rbf', 'sigmoid', 'precomputed'} or callable, default='rbf' Specifies the kernel type to be used in the algorithm. If none is given, 'rbf' will be used. If a callable is given it is used to precompute the kernel matrix.
	degree : int, default=3 Degree of the polynomial kernel function ('poly'). Ignored by all other kernels.
	gamma : {'scale', 'auto'} or float, default='scale' Kernel coefficient for 'rbf', 'poly' and 'sigmoid'. <ul style="list-style-type: none">• if <code>gamma='scale'</code> (default) is passed then it uses $1 / (n_features * X.var())$ as value of gamma,• if 'auto', uses $1 / n_features$.
	<i>Changed in version 0.22:</i> The default value of <code>gamma</code> changed from 'auto' to 'scale'.
	coef0 : float, default=0.0 Independent term in kernel function. It is only significant in 'poly' and 'sigmoid'.
	tol : float, default=1e-3 Tolerance for stopping criterion.
	C : float, default=1.0 Regularization parameter. The strength of the regularization is inversely proportional to C. Must be strictly positive. The penalty is a squared L2 penalty.
	epsilon : float, default=0.1 Epsilon in the epsilon-SVR model. It specifies the epsilon-tube within which no penalty is associated in the training loss function with points predicted within a distance epsilon from the actual value.
	shrinking : bool, default=True Whether to use the shrinking heuristic. See the User Guide .
	cache_size : float, default=200 Specify the size of the kernel cache (in MB).
	verbose : bool, default=False Enable verbose output. Note that this setting takes advantage of a per-process runtime setting in libsvm that, if enabled, may not work properly in a multithreaded context.
	max_iter : int, default=-1 Hard limit on iterations within solver, or -1 for no limit.

Attributes:

- class_weight_ : ndarray of shape (n_classes,)**
Multipliers of parameter C for each class. Computed based on the `class_weight` parameter.
- coef_ : ndarray of shape (1, n_features)**
Weights assigned to the features when `kernel="linear"`.
- dual_coef_ : ndarray of shape (1, n_SV)**
Coefficients of the support vector in the decision function.
- fit_status_ : int**
0 if correctly fitted, 1 otherwise (will raise warning)
- intercept_ : ndarray of shape (1,)**
Constants in decision function.
- n_features_in_ : int**
Number of features seen during `fit`.

New in version 0.24.
- feature_names_in_ : ndarray of shape (n_features_in_,)**
Names of features seen during `fit`. Defined only when `x` has feature names that are all strings.

New in version 1.0.
- n_support_ : ndarray of shape (n_classes,), dtype=int32**
Number of support vectors for each class.
- shape_fit_ : tuple of int of shape (n_dimensions_of_X,)**
Array dimensions of training vector `X`.
- support_ : ndarray of shape (n_SV,)**
Indices of support vectors.
- support_vectors_ : ndarray of shape (n_SV, n_features)**
Support vectors.

Methods

<code>fit(X, y[, sample_weight])</code>	Fit the SVM model according to the given training data.
<code>get_params([deep])</code>	Get parameters for this estimator.
<code>predict(X)</code>	Perform regression on samples in X.
<code>score(X, y[, sample_weight])</code>	Return the coefficient of determination of the prediction.
<code>set_params(**params)</code>	Set the parameters of this estimator.